All You Can Embed: Natural Language based Vehicle Retrieval with Spatio-Temporal Transformers

AI City Challenge 2021

Carmelo Scribano[1,2], Davide Sapienza[1,2], Giorgia Franchini[1,3]
Micaela Verucchi[1] and Marko Bertogna[1]

[1]University of Modena and Reggio Emilia, [2]University of Parma, [3]University of Ferrara

UNIMORE
UNIVERSITÀ DEGLI STUDI DI MODENA E REGGIO EMILIA

High Performance Real Time Lab

UNIVERSITÀ DI PARMA

Università degli Studi di Ferrara

Problem statement

"*Natural language (NL) description offers another useful way to specify vehicle track queries. In this new Challenge Track, participating teams will perform vehicle retrieval given single-camera tracks and corresponding NL descriptions of the targets.*"

- AI City Challenge 2021 – Track5

GOAL:

- Matching Single-Vehicle Tracking Sequences with the corresponding Natural Language Descriptions.

Related tasks:

- Image and Video retrieval
- Multi-Modal video Understanding

1. A white pick up truck drives down the street and passes another truck.
2. A white pickup truck going straight down the street with cars parking on the side.
3. A white pickup runs on the street.

1. A red SUV runs down the street alongside parked cars.
2. Red SUV keeps straight followed by a maroon car.
3. A red SUV runs down the road followed by a black vehicle.

Model Overview

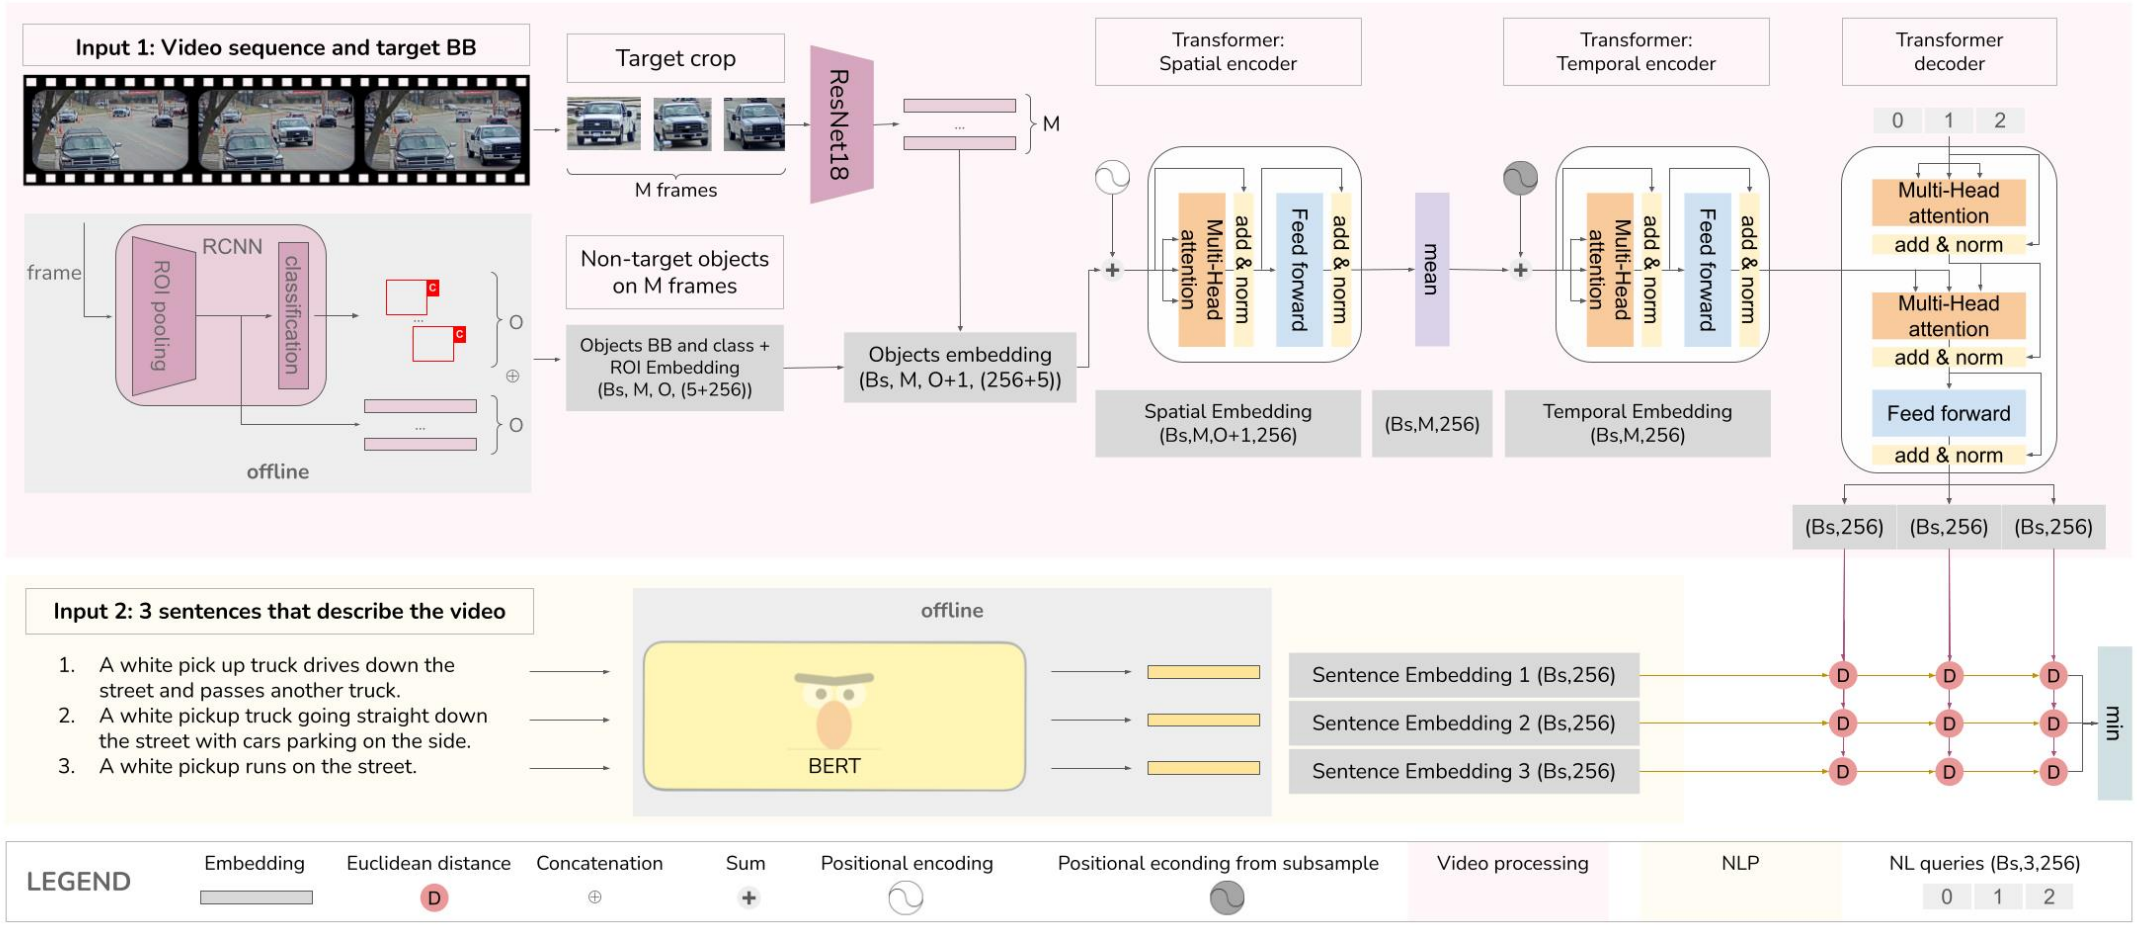

Natural Language Branch

OBJECTIVE:

$$\begin{cases} d\left(T_1^i, T_2^i\right) \ll d\left(T_1^i, T_2^j\right) \\ d\left(T_1^i, T_2^i\right) \approx 0 \end{cases} \quad \forall i : i \neq j$$

Where:

$$T_a^i = BERT\left(t_a^i\right) \qquad d(u,v) = 1 + \frac{u^T v}{||u||||v||} \in [0,2]$$

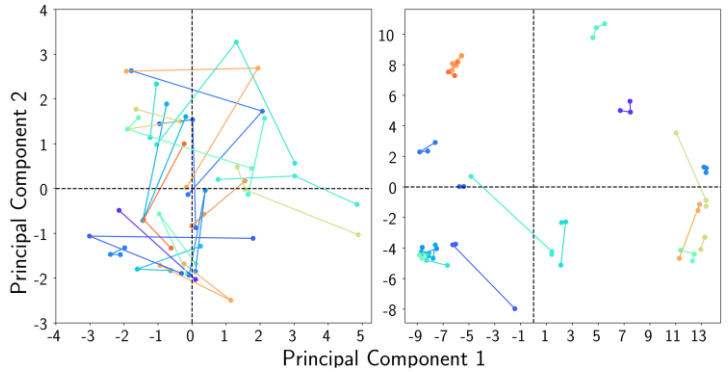

BEFORE:

$$mean(d_{INTER}) = 0.1703$$

$$mean(d_{INTRA}) = 0.1899$$

AFTER:

$$mean(d_{INTER}) = \mathbf{0.2089}$$

$$mean(d_{INTRA}) = \mathbf{0.6140}$$

POSITIVE
"A grey sedan turning right."

ANCHOR
"A silver car is driving in the right lane around a curve"

NEGATIVE
"A red SUV runs down the street alongside parked cars."

Randomly chosen (A,P,N) Triple

BERT (BASE)

Pre-Trained BERT[1] Model

768-dimensional Output

$$TL = \frac{1}{Bs}\sum_{i=1}^{Bs} \max(0, d\left(A^i, P^i\right) - d\left(A^i, N^i\right) + m)$$

Triplet Loss

[1]: Bert:Pre-training of deep bidirectional transformers for language understanding

Visual Branch

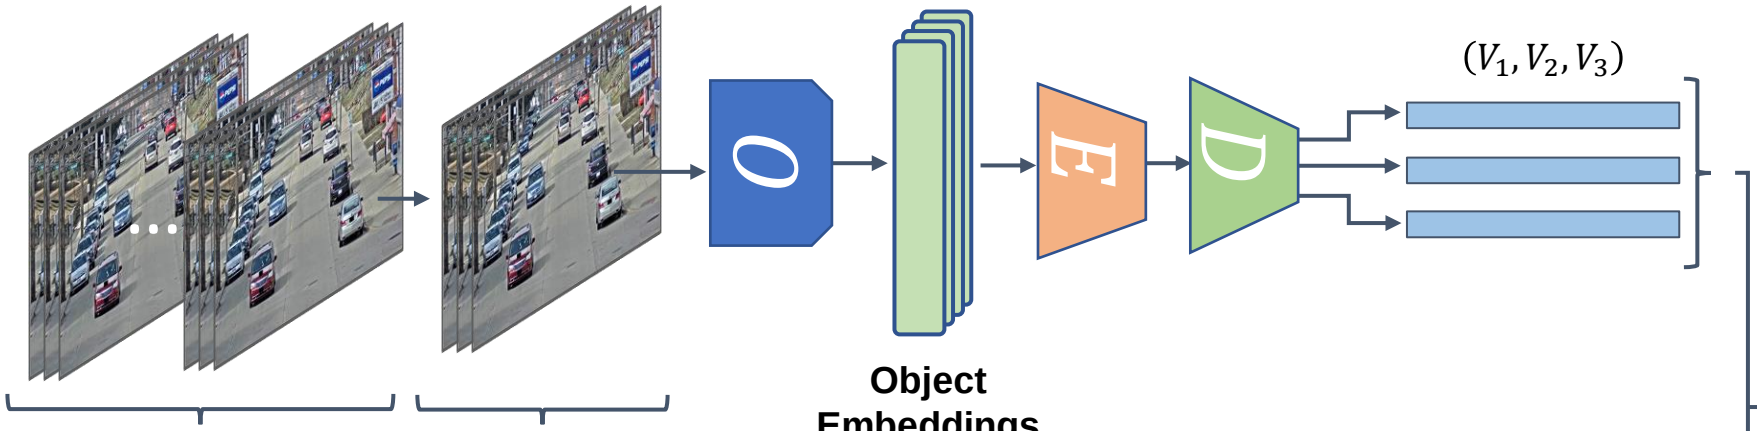

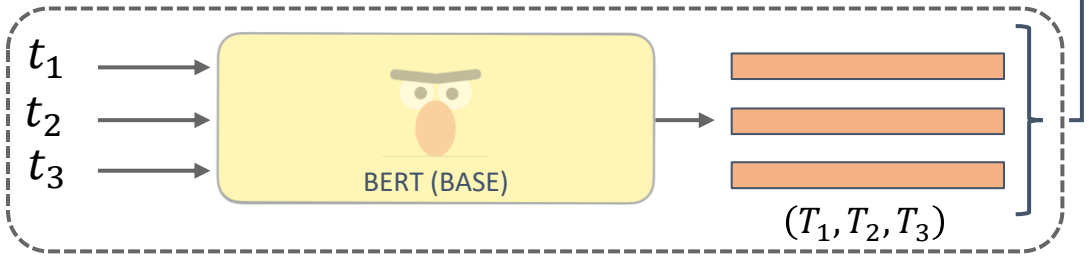

Language Branch
*After Finetuning

Evaluation Mode (Inference)

$$d(V,T) = \min\left(d(V_i, T_j)\right) \quad \forall\ (i,j)$$

Object Embeddings

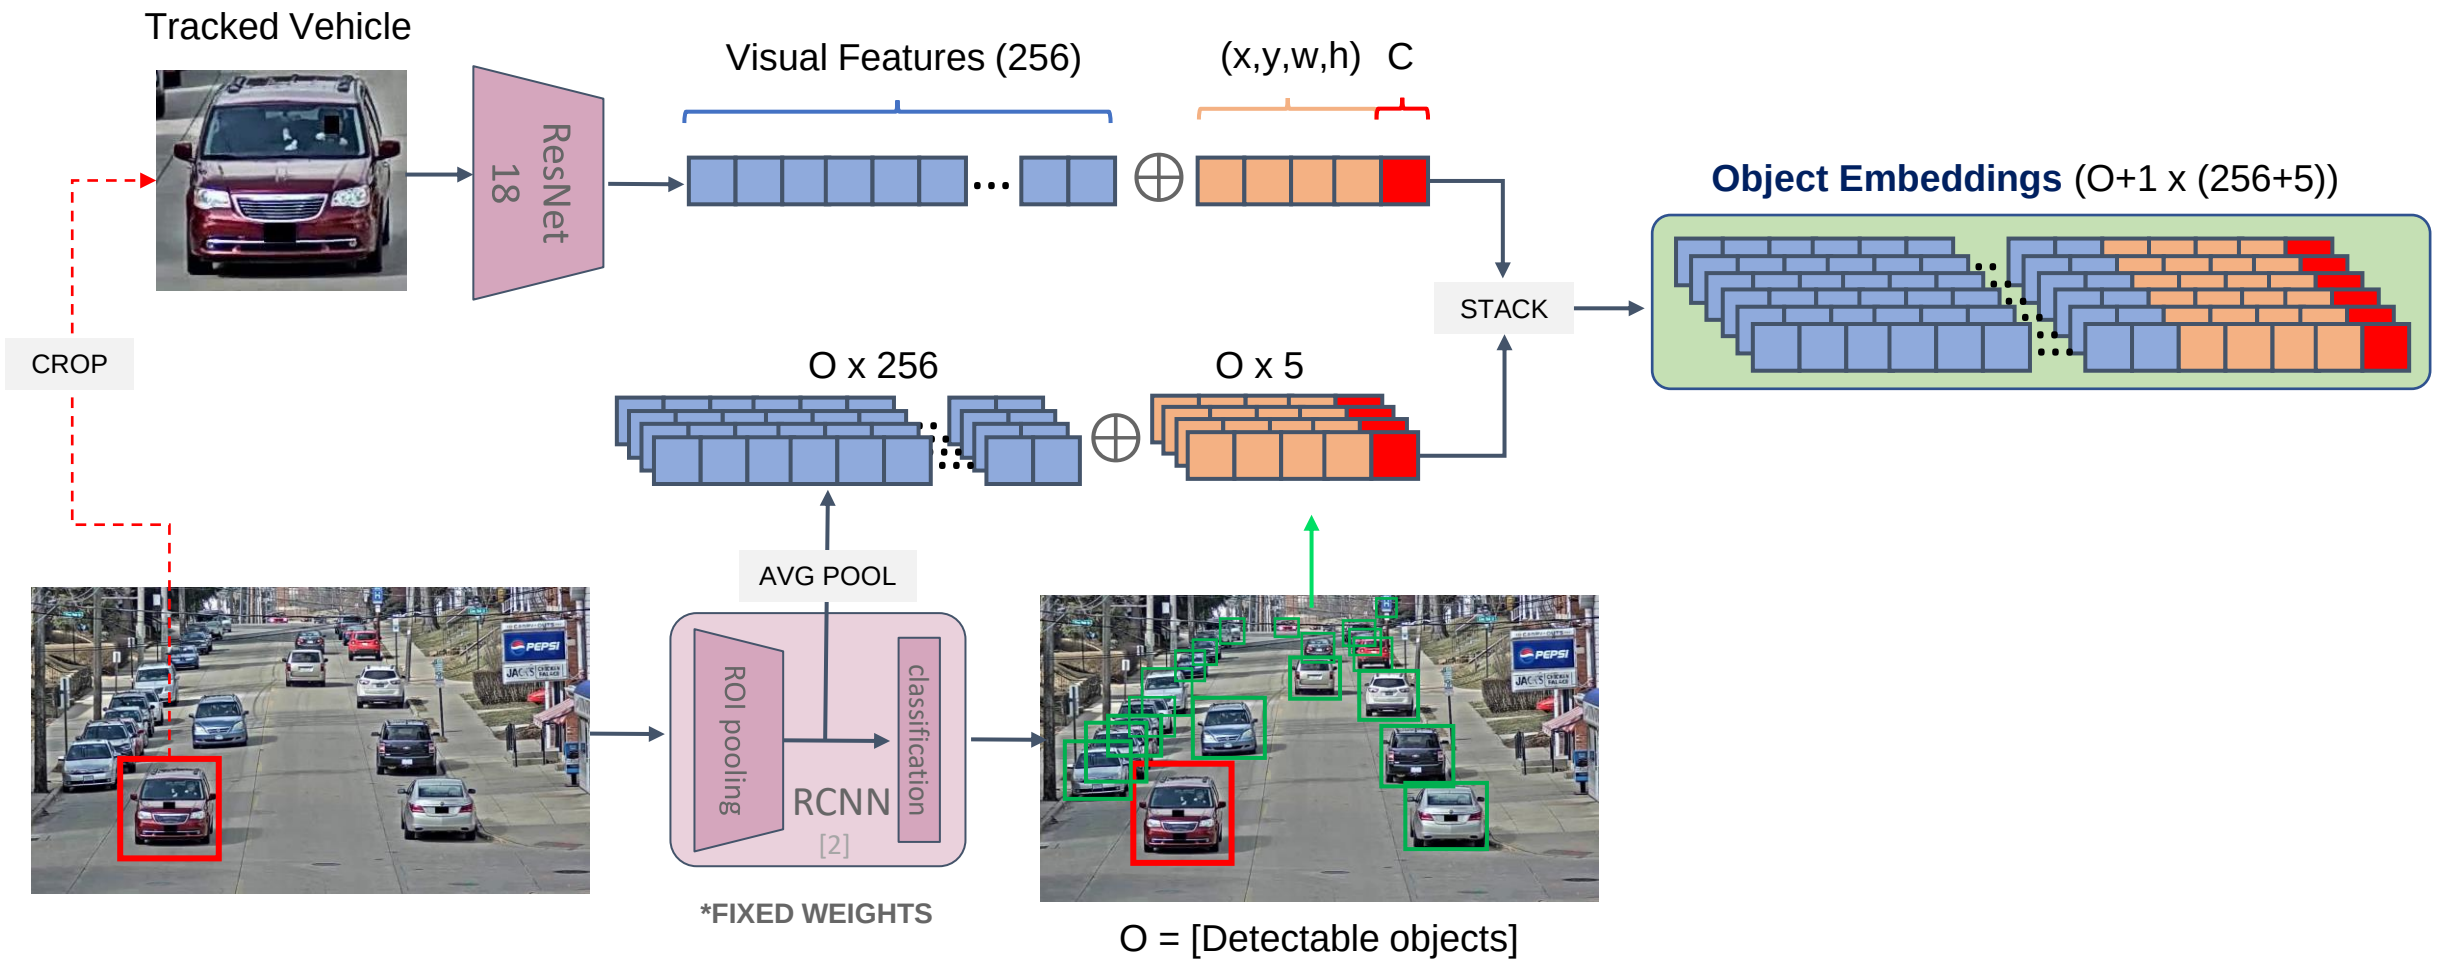

[2]: Faster r-cnn: Towards real-time object detection with region proposal networks

Spatio-Temporal Transformers

RATIONALE:

- The Spatial Encoder provides a comprehensive frame-level representation
- The Temporal Encoder encodes the sequential meaning of the depicted scene.

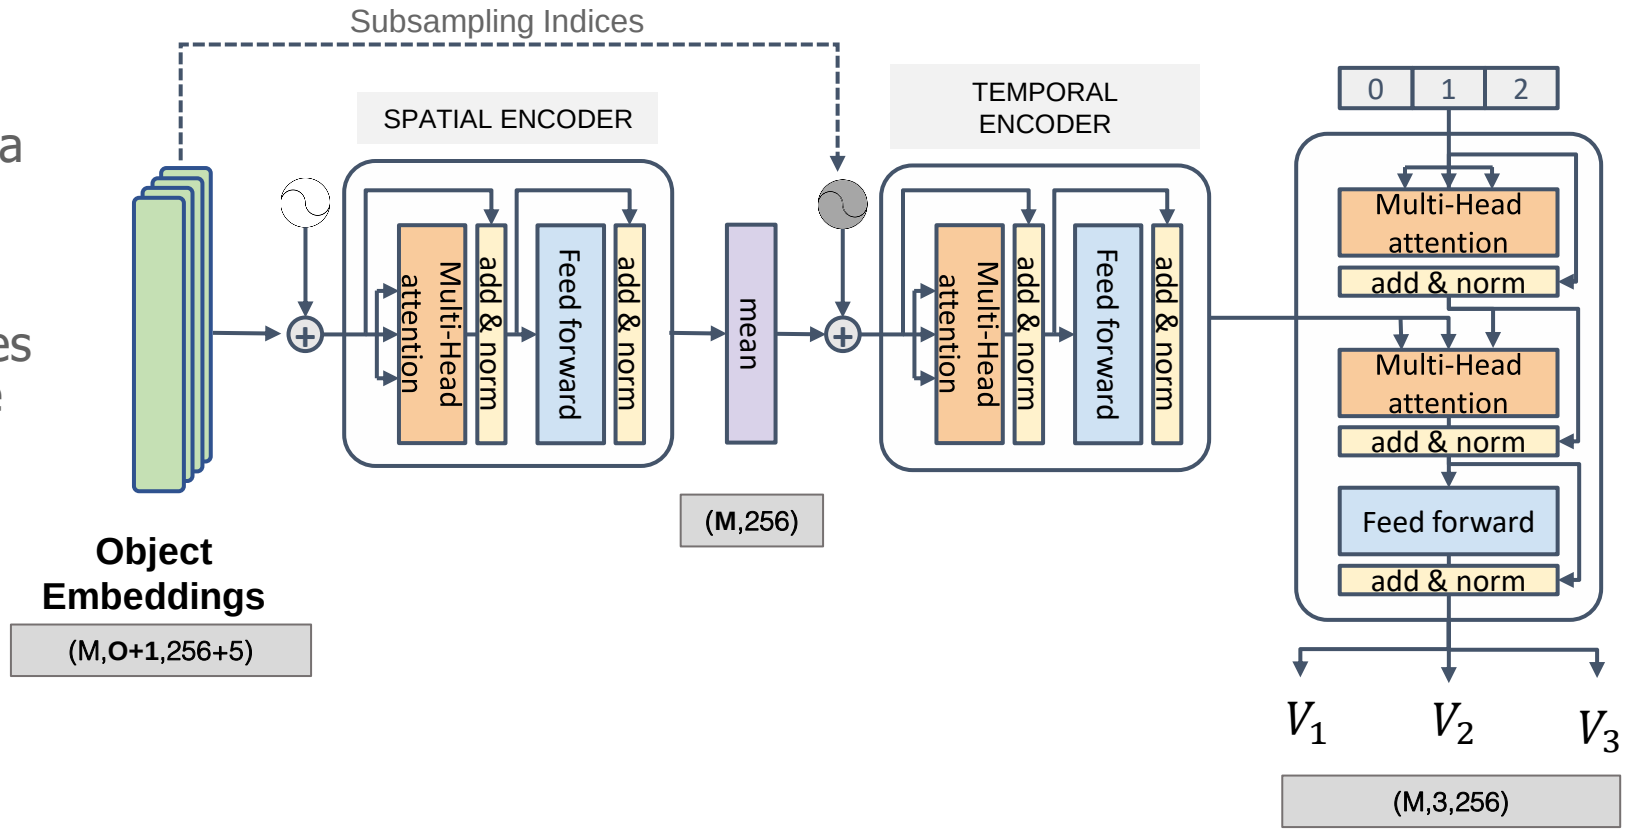

[3]: Attention is All you Need

Optimization

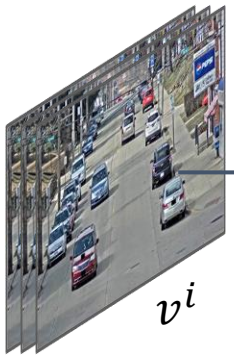

v^i

Visual Branch

ANCHOR

$A_n^i \quad n \in \{1,2,3\}$

*FIXED WEIGHTS

$t^i_{\{1,2,3\}}$

$t^j_{\{1,2,3\}}$

Language Branch

$i \neq j$

POSITIVE

$P_n^i \quad n \in \{1,2,3\}$

$N_n^i \quad n \in \{1,2,3\}$

NEGATIVE

$$\mathcal{L}_{AYCE}(A,P,N) = TL(A,P,N) + \frac{1}{Bs}\sum_{i=1}^{Bs} \beta \cdot \Phi\left(A^i, P^i\right)$$

$$\begin{cases} \Phi(A,P) = \min\left(d\left(A_{m,}P_n\right)\right) \quad m,n \in \{1,2,3\} \\ d_{TL}(A,[P|N]) = \frac{1}{9}\sum_{m=1}^{3}\sum_{n=1}^{3}\|A_m - [P|N]_n\|^2 \end{cases}$$

RESULTS

MODEL	MRR	RE@5	RE@10
BASELINE	0.0269	0.0264	0.0491
OURS (BEST – 11°)	**0.1078**	0.1321	0.2491
Alibaba-UTS-ZJU (*1°*)	0.1869	-	-
Sun Asterisk (4°)	0.1571	-	-
Modulab (10°)	0. 1195	-	-

Qualitative Results

"A red pickup drives straight down a *highway.“*

"A red pickup truck runs down the street.“

"A red pickup following straight other three car”



"A blue pickup runs down the street."

"A blue pickup truck going straight down the street passing an intersection."

"A blue truck runs down the street.“



"A black sedan crossing an intersection."

"A midsize black sedan goes straight through the intersection."

"A black Sedan runs down the street."